

Business Case: Bicycle Sales Data Warehouse



Introduction

The data warehouse for Bicycle Sales aims to consolidate sales, customer, inventory, and staff data into a unified platform for comprehensive analysis and decision-making. By integrating various data sources, the company will be able to optimize sales, manage inventory effectively, and enhance customer experience. This report discusses the business case, key business questions, the data warehouse answers, the rationale behind its schema design, and the data sources and integration methods employed.

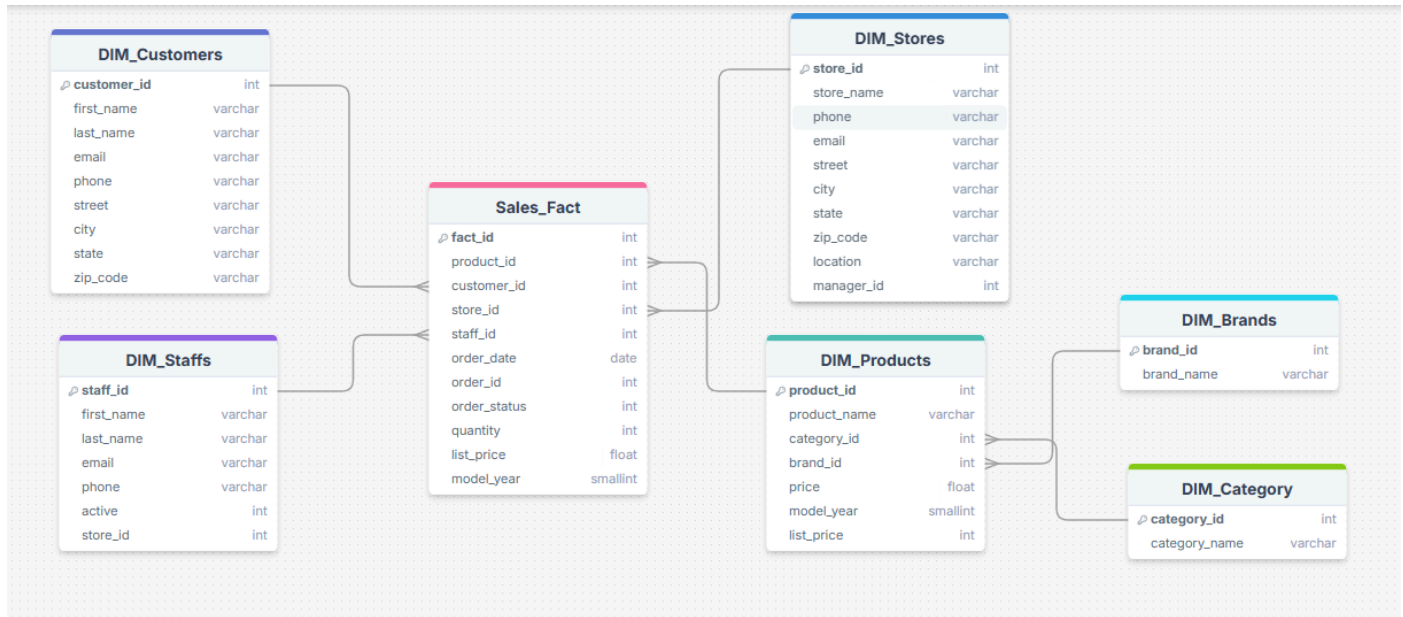
The Business Case for the Data Warehouse

The bicycle sales company operates across multiple sales channels, offering various types of bicycles (e.g., road bikes, mountain bikes, and e-bikes). With the increasing volume of sales transactions, customer interactions, and product offerings, the company faces challenges in managing data from disparate systems. These challenges include limited visibility into sales performance, inventory management issues, and the inability to target customers effectively with personalized marketing strategies.

A data warehouse solves these problems by centralizing the data from multiple sources, such as sales transactions, customer databases, and inventory systems, into one cohesive platform. The data warehouse will enable better decision-making, streamline reporting, improve forecasting, and help the company respond more effectively to market trends. The integration of data from multiple channels will allow the company to understand sales patterns, optimize inventory, and increase revenue by leveraging data insights.

Snowflake Schema design :

The schema design of the bicycle sales data warehouse is structured around a Snowflake schema, which is ideal for analytical queries and decision-making. The fact table ([Sales_Fact](#)) acts as the central hub, linking to various dimension tables ([Products](#), [Customers](#), [Stores](#), [Staffs](#), and [Categories](#)). This design facilitates efficient querying, as fact tables store transactional data, and dimension tables store descriptive attributes.



Key features of the schema design:

- **Fact Table (Sales_Fact):** The fact table stores the core transactional data, including sales transactions, order dates, quantities, product details, store details, and staff involved in the sale. This table is connected to dimension tables through foreign keys.
- **Dimension Tables:** These tables provide detailed context for the sales data:
 - **Products Dimension:** Contains product details, including name, price, category, and brand.
 - **Customers Dimension:** Stores customer demographic information, enabling customer segmentation.
 - **Stores Dimension:** Details of each store, including location and management.
 - **Staffs Dimension:** Contains staff details to analyze performance and track sales by staff members.
 - **Categories and Brands Dimensions:** Used to group products by category (e.g., road bikes, mountain bikes) and brand, enabling detailed performance analysis.

Inside the zip file:

- Raw data files in the attached folder
- Python file that loads data into PostgreSQL.
- Files to create original tables and snowflake schema tables.
- Answers to business questions.

Business Questions Answered

The data warehouse answers key business questions that drive operational improvements and growth. These questions include:

- **Which bicycle categories are generating the most revenue?** The warehouse aggregates sales data by product category (e.g., road bikes, mountain bikes) and provides the total revenue generated by each category. This insight helps the company understand which product lines are most profitable and decide where to focus marketing and sales efforts.

The screenshot shows a PostgreSQL query editor with a query to find the top bicycle categories by revenue. The query is as follows:

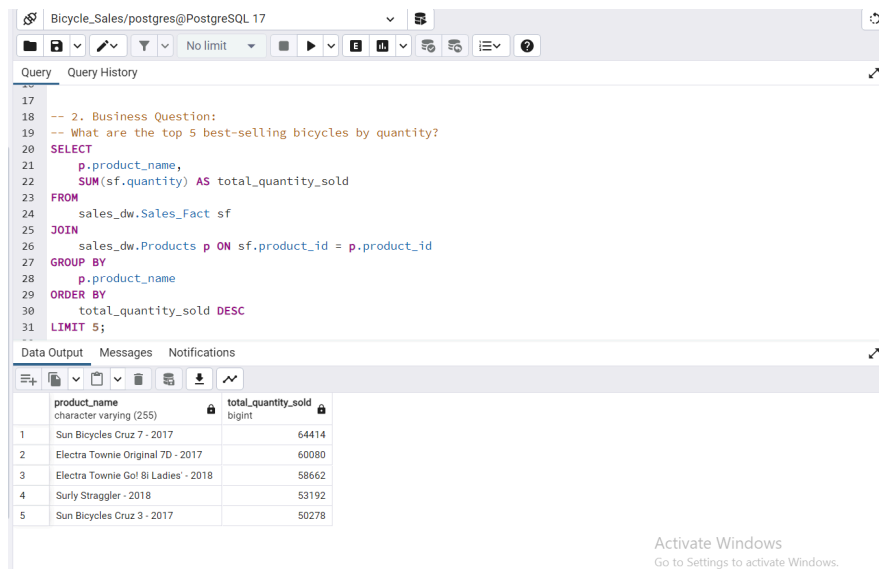
```
1 -- 1. Business Question:
2 -- Which bicycle categories (e.g., road bikes, mountain bikes, e-bikes) are generating the most revenue?
3 SELECT
4     c.category_name,
5     SUM(sf.quantity * p.list_price) AS total_revenue
6 FROM
7     sales_dw.Sales_Fact sf
8 JOIN
9     sales_dw.Products p ON sf.product_id = p.product_id
10 JOIN
11     sales_dw.Categories c ON p.category_id = c.category_id
12 GROUP BY
13     c.category_name
14 ORDER BY
15     total_revenue DESC;
```

The results are displayed in a table with the following columns: category_name (character varying (255)) and total_revenue (bigint). The results are ordered by total_revenue in descending order.

	category_name	total_revenue
1	Road Bikes	3442756232
2	Mountain Bikes	1944352768
3	Electric Bikes	1667992880
4	Cruisers Bicycles	1198965716
5	Cyclocross Bicycles	403188740
6	Children Bicycles	334094780
7	Comfort Bicycles	132804076

- **What are the top-selling bicycles by quantity?**

- The data warehouse allows the company to identify the top-selling bicycles by quantity over a specific time period. This helps with stock planning, ensuring that popular bikes are always in stock while underperforming models can be phased out or promoted.



The screenshot shows a PostgreSQL query editor with the following SQL query:

```

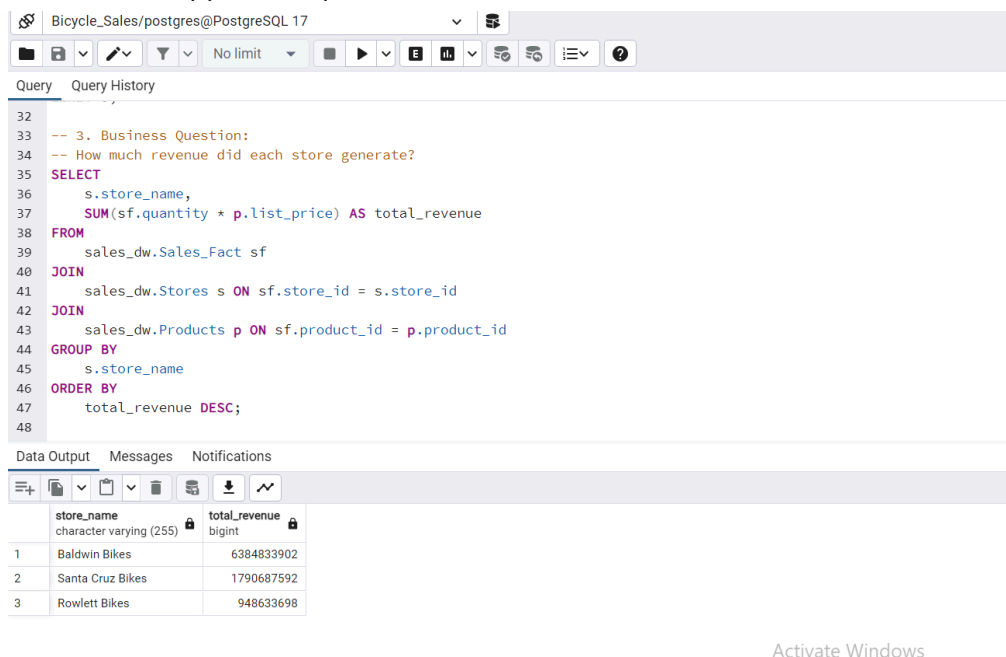
17
18 -- 2. Business Question:
19 -- What are the top 5 best-selling bicycles by quantity?
20 SELECT
21     p.product_name,
22     SUM(sf.quantity) AS total_quantity_sold
23 FROM
24     sales_dw.Sales_Fact sf
25 JOIN
26     sales_dw.Products p ON sf.product_id = p.product_id
27 GROUP BY
28     p.product_name
29 ORDER BY
30     total_quantity_sold DESC
31 LIMIT 5;

```

The query results are displayed in a table with the following data:

	product_name	total_quantity_sold
1	Sun Bicycles Cruz 7 - 2017	64414
2	Electra Townie Original 7D - 2017	60080
3	Electra Townie Go! 8i Ladies - 2018	58662
4	Surly Straggler - 2018	53192
5	Sun Bicycles Cruz 3 - 2017	50278

- **How much revenue is generated by each store?** By linking sales data with store locations, the data warehouse provides detailed performance metrics at the store level. This information helps identify high-performing stores and areas that require additional support or promotional efforts.



The screenshot shows a PostgreSQL query editor with the following SQL query:

```

32
33 -- 3. Business Question:
34 -- How much revenue did each store generate?
35 SELECT
36     s.store_name,
37     SUM(sf.quantity * p.list_price) AS total_revenue
38 FROM
39     sales_dw.Sales_Fact sf
40 JOIN
41     sales_dw.Stores s ON sf.store_id = s.store_id
42 JOIN
43     sales_dw.Products p ON sf.product_id = p.product_id
44 GROUP BY
45     s.store_name
46 ORDER BY
47     total_revenue DESC;
48

```

The query results are displayed in a table with the following data:

	store_name	total_revenue
1	Baldwin Bikes	6384833902
2	Santa Cruz Bikes	1790687592
3	Rowlett Bikes	948633698